

```

import pandas as pd
import seaborn as sns
import tensorflow as tf
from tensorflow import keras
import numpy as np

# Load dataset. This dataset is available at https://www.kaggle.com/datasets/datasnaek/mbti
# The full details of the dataset are available on Kaggle.
# The dataset contains a Myers-Briggs Personality type, which will serve as the target,
# and a collection of parts of their 50 most recent posts, which serve as the text input.
# I will split up the text and maintain the personality type to make the model more general
# applicable.
df = pd.read_csv('mbti_1.csv', dtype={'posts':str, 'type':str})

```

```
df.head()
```

| | type | posts |
|---|------|---|
| 0 | INFJ | 'http://www.youtube.com/watch?v=qsXHcwe3krw ... |
| 1 | ENTP | 'I'm finding the lack of me in these posts ver... |
| 2 | INTP | 'Good one _____ https://www.youtube.com/wat... |
| 3 | INTJ | 'Dear INTP, I enjoyed our conversation the o... |
| 4 | ENTJ | 'You're fired. That's another silly misconce... |

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8675 entries, 0 to 8674
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    type    8675 non-null        object
1    posts   8675 non-null        object
dtypes: object(2)
memory usage: 135.7+ KB

```

```
df['type'].unique()
```

```

array(['INFJ', 'ENTP', 'INTP', 'INTJ', 'ENTJ', 'ENFJ', 'INFP', 'ENFP',
       'ISFP', 'ISTP', 'ISFJ', 'ISTJ', 'ESTP', 'ESFP', 'ESTJ', 'ESFJ'],
      dtype=object)

```

```

# Some personality types are more well-represented than others
df['type'].value_counts()

```

```

INFP    1832
INFJ    1470
INTP    1304
INTJ    1091
ENTP    685
ENFP    675
ISTP    337
ISFP    271
ENTJ    231
ISTJ    205
ENFJ    190
ISFJ    166
ESTP    89
ESFP    48
ESFJ    42
ESTJ    39
Name: type, dtype: int64

```

```


# Graph the distribution of target classes (equal amounts of Kecimen and Besni species)
sns.countplot(x=df['type'])

```

```

<AxesSubplot: xlabel='type', ylabel='count'>

```



```

# Convert compiled texts into lists
df['text'] = df['posts'].apply(lambda e: e.split('|||')) # for now just remove the |||
df.drop('posts', axis=1, inplace=True)
df['text'].head()

```

```

0    ['http://www.youtube.com/watch?v=qsXHcwe3krw, ...
1    ['I'm finding the lack of me in these posts ve...
2    ['Good one _____ https://www.youtube.com/wa...
3    ['Dear INTP, I enjoyed our conversation the ...
4    ['You're fired., That's another silly misconce...
Name: text, dtype: object

```

```

# Convert lists into individual rows; 1 row -> 5 rows, and stop associating them by ignorir
df = df.explode('text', ignore_index=True)
df.head()

```

| | type | text |
|---|------|---|
| 0 | INFJ | 'http://www.youtube.com/watch?v=qsXHcwe3krw |
| 1 | INFJ | http://41.media.tumblr.com/tumblr_lfouy03PMA1q... |
| 2 | INFJ | enfp and intj moments https://www.youtube.com... |
| 3 | INFJ | What has been the most life-changing experienc... |
| 4 | INFJ | http://www.youtube.com/watch?v=vXZeYwwRDw8 h... |

```

# The dataset is enormous now
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 422845 entries, 0 to 422844
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    type    422845 non-null    object
1    text    422845 non-null    object
dtypes: object(2)
memory usage: 6.5+ MB

```

```

# Let's trim it down.
df = df.truncate(after=10000)

```

```

# Split the DataFrame into features and target
X = df['text']
y = df['type']

```

```

# Create a dictionary to convert classes to numbers for the model
enumerated_classes = {key: i for i, key in enumerate(y.unique())}
enumerated_classes

```

```

{'TNF1': 0.

```

```
{ 'ENTP': 1,  
  'INTP': 2,  
  'INTJ': 3,  
  'ENTJ': 4,  
  'ENFJ': 5,  
  'INFP': 6,  
  'ENFP': 7,  
  'ISFP': 8,  
  'ISTP': 9,  
  'ISFJ': 10,  
  'ISTJ': 11,  
  'ESTP': 12,  
  'ESFP': 13}
```

```
# Convert y to a list of numeric values
```

```
y = y.apply(lambda type_name: enumerated_classes[type_name]  
            )
```

```
# Divide the data into train and test
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
# Ensure we chose the right variables from the train test split
```

```
len(X_train) == len(y_train) and len(X_test) == len(y_test)
```

```
True
```

```
# Vectorize the text
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
tokenizer = Tokenizer(num_words=400) # Only get top 400 terms to conserve memory
```

```
tokenizer.fit_on_texts(X_train)
```

```
# Vectorized data
```

```
X_train_vectorized = tokenizer.texts_to_sequences(X_train)
```

```
X_test_vectorized = tokenizer.texts_to_sequences(X_test)
```

```
X_train_vectorized
```

```
[[223, 16, 40, 160, 13, 1, 226, 265, 5, 24, 26, 1, 287, 57],  
 [2, 371, 15, 55, 33, 1, 78, 43, 2, 314, 1, 287, 4, 14, 198, 61, 1, 3, 2],  
 [1,  
  1,  
  124,  
  14,  
  5,  
  195,  
  61,  
  256,  
  ^
```

```
3,  
244,  
124,  
198,  
191,  
90,  
1,  
32,  
339,  
3,  
67,  
21,  
1,  
219,  
126,  
3,  
244,  
124,  
164,  
5,  
158,  
90,  
44,  
191,  
9,  
353,  
1],  
[1, 36, 51, 22, 51, 42, 29, 47, 82, 99, 74, 1, 223, 16, 40, 216],  
[19, 3, 2, 39, 26, 162],  
[],  
[1,  
31,  
280,  
4,  
18,  
232,  
8,  
2,  
34,  
40,  
60,  
8,  
59,  
207,  
20,  
323,  
4,  
74,
```

```
# Pad the sequences so that they are the same length  
from tensorflow.keras.preprocessing.sequence import pad_sequences  
X_train_vectorized = pad_sequences(X_train_vectorized, maxlen=50, padding='post', truncating='post')  
X_test_vectorized = pad_sequences(X_test_vectorized, maxlen=50, padding='post', truncating='post')
```

```
X_train_vectorized
```

```
array([[223., 16., 40., ..., 0., 0., 0.],
       [ 2., 371., 15., ..., 0., 0., 0.],
       [ 1., 1., 124., ..., 0., 0., 0.],
       ...,
       [341., 51., 22., ..., 0., 0., 0.],
       [ 33., 29., 7., ..., 0., 0., 0.],
       [ 1., 398., 106., ..., 0., 0., 0.]], dtype=float32)
```

```
seq_model = keras.Sequential([
    keras.layers.Embedding(400, 32, input_length=50),
    keras.layers.LSTM(64, dropout=0.1),
    keras.layers.Dense(len(enumerated_classes), activation='softmax')
])
```

```
seq_model.summary()
```

```
Model: "sequential"
```

| Layer (type) | Output Shape | Param # |
|--------------------------|----------------|---------|
| ===== | | |
| embedding (Embedding) | (None, 50, 32) | 12800 |
| lstm (LSTM) | (None, 64) | 24832 |
| dense (Dense) | (None, 14) | 910 |
| ===== | | |
| Total params: 38,542 | | |
| Trainable params: 38,542 | | |
| Non-trainable params: 0 | | |

```
seq_model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)
```

```
# seq_model.fit(X_train_vectorized, y_train, epochs=5)
# My PC completely crashes and restarts a couple seconds after running this so... I won't.
# Instead I will continue with the rest of the code but will not be able to run it.
```

```
Epoch 1/5
250/250 [=====] - 2s 7ms/step - loss: 2.2231 - accuracy: 0.1
Epoch 2/5
244/250 [=====>.] - ETA: 0s - loss: 2.2221 - accuracy: 0.1820
```

```
# Check for overfitting by scoring against test values
```

```
seq_model.evaluate(X_test_vectorized, y_test)
```

```
63/63 [=====] - 0s 4ms/step - loss: 2.2522 - accuracy: 0.170  
[2.2521889209747314, 0.17091454565525055]
```

```
# Create a RNN
```

```
rnn = keras.Sequential([  
    keras.layers.Embedding(input_dim=400, output_dim=60),  
    keras.layers.LSTM(100),  
    keras.layers.Dense(len(enumerated_classes)) # Output layer  
)
```

```
rnn.summary()
```

```
Model: "sequential_2"
```

| Layer (type) | Output Shape | Param # |
|-------------------------|------------------|---------|
| embedding_2 (Embedding) | (None, None, 60) | 24000 |
| lstm_2 (LSTM) | (None, 100) | 64400 |
| dense_2 (Dense) | (None, 14) | 1414 |

```
=====
```

```
Total params: 89,814  
Trainable params: 89,814  
Non-trainable params: 0
```

```
=====
```

```
rnn.evaluate(X_test_vectorized, y_test)
```

```
# Trying different embedding approaches
```

```
# Instead of assigning each word an integer and listing them in sequence, assign each sentence  
from sklearn.feature_extraction.text import CountVectorizer
```

```
vectorizer = CountVectorizer(lowercase=False, max_df=0.90, min_df=0.01) # Ignore terms that
```

```
vectorizer.fit(X_train)
```

```
# Vectorized data
```

```
X_train_vectorized = vectorizer.transform(X_train).toarray()
```

```
X_test_vectorized = vectorizer.transform(X_test).toarray()
```

```
X_train_vectorized = pad_sequences(X_train_vectorized, maxlen=50, padding='post', truncating='post')
```

```
X_test_vectorized = pad_sequences(X_test_vectorized, maxlen=50, padding='post', truncating='post')
```

```
seq_model = keras.Sequential([  
    keras.layers.Embedding(400, 32, input_length=50),
```

```
keras.layers.LSTM(64, dropout=0.1),
keras.layers.Dense(len(enumerated_classes), activation='softmax')
])

seq_model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

seq_model.fit(X_train_vectorized, y_train, epochs=5)

seq_model.evaluate(X_test_vectorized, y_test)

63/63 [=====] - 1s 3ms/step - loss: 2.2475 - accuracy: 0.193
[2.247462034225464, 0.19340330362319946]
```

The tokenizing approach where each (top) word is assigned an integer value is likely to perform much better than the bag-of-words style vectorization because it encodes within it the sequential data of the text that is lost when doing a bag-of-words. Unfortunately I could not run all the models to verify the difference in performance due to my poor setup or whatever the cause of my computer crashing is if not that. Other existing pre-trained models like BERT could serve as alternative embeddings which would likely improve the performance of this text classifier.

• •